

SPREADSHEETLLM Framework

- **Goal:** Enable LLMs to “understand” and reason over full spreadsheets despite token-length and 2D-layout challenges.
- **Initial (Vanilla) Encoding:**
 - Serialize every cell as a sequence including its address (e.g. “A1”), value, and format.
 - *Limitation:* Large sheets exceed model token limits and performance degrades.

- **SHEETCOMPRESSOR Encoding Pipeline**

To dramatically shrink input size while preserving layout and content, three modular steps are applied in sequence:

1. **Structural-Anchor Compression**
 - Detect “anchors” (rows/columns where content changes) that mark table boundaries.
 - Discard distant, homogeneous rows/columns to yield a compact “skeleton” of key structure.
 2. **Inverted-Index Translation**
 - Build a JSON dictionary mapping each unique non-empty cell text to all its addresses.
 - Replace repeated values with index lookups, preserving losslessness but slashing tokens.
 3. **Data-Format-Aware Aggregation**
 - Group adjacent numeric cells sharing the same format string (e.g. “\$#,##0.00”) or data type.
 - Represent each rectangular block by its format, avoiding redundant numeric 1.
- Structural-Anchor-Based Extraction**

Goal: Identify and retain only structurally important (heterogeneous) rows/columns to produce a skeleton version of the spreadsheet.

Implementation Steps:

- **Input:** Raw spreadsheet with all cells.
- **Step 1 – Detect Structural Anchors:**
 - Use heuristics or statistical features to identify candidate table boundaries (e.g., rows with headers, or columns that differ significantly in format/content from neighbors).
 - These are the yellow-highlighted vertical and horizontal anchor lines in panels (a) and (b).
- **Step 2 – Remove Redundant Rows/Cols:**

- Delete rows/columns that are homogeneous and far from detected anchors (e.g., repetitive numeric rows far from headers).
 - Retain only surrounding context (within k cells) from each structural anchor.
 - **Output:** A significantly reduced spreadsheet—e.g., from $576 \times 23 \rightarrow 24 \times 8$ (as in panel (c)).
-

2. Inverted-Index Translation

Goal: Optimize serialization by grouping identical cell values and storing them as key-value pairs where the key is the value and the value is a list of cell addresses.

Implementation Steps:

- **Input:** Skeleton sheet from Step 1.
- **Step 1 – Traverse the Sheet:**
 - For each non-empty cell, check if its value already exists in a dictionary.
 - If so, append the cell's coordinate (e.g., B4) to that value's list. If not, create a new entry.
- **Step 2 – Represent as JSON Dictionary:**
 - Structure: { "cell_text": [list of coordinates] }
 - Examples from the encoding block:

```
json
CopyEdit
{
  "Atlantis": ["A2", "A7", "F5"],
  "QuantumMind": ["B2", "B3", "B4", "F4"],
  "20-Aug": ["D2:D18", "D21:D23"]
}
```

- **Output:** A JSON index that eliminates repeated values and avoids encoding empty cells.
-

3. Data-Format-Aware Aggregation

Goal: Group and represent adjacent numeric cells that share identical formats to further compress the spreadsheet.

Implementation Steps:

- **Input:** Inverted-index-translated sheet.
- **Step 1 – Scan for Format Similarities:**
 - Identify contiguous numeric cells (same data type or format string).
 - Examples: currency (\$#, ##0.00), integer, percentage.

- **Step 2 – Group Cells by Format:**
 - Merge them into labeled format clusters, e.g.:

```
json
CopyEdit
{
  "IntNum": ["C2:C4", "C6:C12"],
  "Percentage": ["H5:H6", "H8:H9"]
}
```

- **Output:** Further compressed index that encodes ranges by shared formatting blocks.

3 Method: Overview

We represent a sheet S as an $m \times n$ grid of cells and turn it into text via three independent, composable modules. First up:

3.1 Vanilla Spreadsheet Encoding

Goal: Turn every cell into a short Markdown-style line containing its address, value, and (optionally) format.

1. Data Structures

```
# S: 2D array or dict mapping (i,j) → Cell(value, format)
text_lines: List[str] = []
```

2. Loop over all cells

```
for i in range(1, m+1):
    for j in range(1, n+1):
        cell = S.get((i,j))
        if not cell or cell.is_empty():
            continue
        addr = f"{col_letter(j)}{i}"          # e.g. A1, B3, ...
        # Omit `cell.format` if you want faster token usage.
        line = f"|{addr}|{cell.value}|{cell.format}|"
        text_lines.append(line)
```

3. Concatenate into a single string

```
T_vanilla = "\n".join(text_lines)
```

Note: Including rich format metadata (colors, borders, fonts) rapidly exhausts LLM token limits and often hurts accuracy—so default to value + address only, or prune formats to the essentials.

3.2 Structural-Anchor-Based Extraction

Goal: Strip away “boring” homogeneous rows/columns far from any table boundary, keeping only a small “skeleton” that preserves layout cues.

1. Measure Heterogeneity

- For each row p , compute $\text{row_score}[p]$ = number of distinct cell values or format changes in that row.
- For each column q , compute $\text{col_score}[q]$ similarly.

2. Select Anchors

```
anchors_rows = { p | row_score[p] ≥  $\theta_{\text{row}}$  }  
anchors_cols = { q | col_score[q] ≥  $\theta_{\text{col}}$  }  
A = anchors_rows  $\cup$  anchors_cols
```

You can also include first/last row+col by default.

3. Define Neighborhood Radius

```
CopyEdit  
k = user_defined_threshold # e.g. 2 or 3
```

4. Build Retained Row/Col Sets

```
kept_rows =  $\cup_{\{p \in \text{anchors\_rows}\}} \{ i \mid |i - p| \leq k \}$   
kept_cols =  $\cup_{\{q \in \text{anchors\_cols}\}} \{ j \mid |j - q| \leq k \}$ 
```

5. Extract Skeleton Cells

```
CopyEdit  
S_skeleton = { (i,j): S[i,j]  
               for i in kept_rows  
               for j in kept_cols  
               if S[i,j] not empty }
```

6. Coordinate Re-mapping

```
# Remap old row indices to 1...m', and cols to 1...n'  
new_row_map = { old_i: new_i for new_i, old_i in  
                enumerate(sorted(kept_rows), start=1) }  
new_col_map = { old_j: new_j for new_j in enumerate(sorted(kept_cols),  
                start=1) }  
  
S_extracted = {  
    (new_row_map[i], new_col_map[j]): cell  
    for (i,j), cell in S_skeleton.items()  
}
```

This preserves the logical relationships (e.g. “A1” → “A1” in the compressed grid).

7. Compression Stats

- Typically drops ~ 75% of cells
- Still retains ~ 97% of boundary rows/columns

8. Output for Next Step

```
Se = S_extracted
Te = vanilla_encode(Se) # reuse 3.1 on the smaller S_extracted
```

3.3 Inverted-Index Translation

Goal: Turn a flat list of (address, value) pairs into a **lossless** dictionary { value → [address or address_range, ...] }, dropping empties and merging repeats.

◆ Implementation Steps

1. Gather Non-Empty Cells

```
# Input: Se from §3.2
cells = [ (addr, cell.value)
          for addr, cell in Se.items()
          if cell.value not in ("", None) ]
```

2. Build Value→Address Map

```
from collections import defaultdict
index = defaultdict(list)
for addr, val in cells:
    index[val].append(addr) # e.g. "A4", "B2", ...
```

3. Compress Addresses into Ranges

- **Sort** each list of addresses.
- **Group** contiguous runs in the same row or column into compact ranges (e.g. "C2", "C3", "C4" → "C2:C4"; "D2", "E2", "F2" → "D2:F2").
- **Function Sketch:**

```
def compress_runs(addr: List[str]) -> List[str]:
    # parse each addr into (row, col), cluster contiguous,
    # then re-serialize as "A1" or "A1:A3"
    ...
```

```
for val, addr_list in index.items():
    index[val] = compress_runs(sorted(addr_list))
```

5. Output as JSON

```
python
CopyEdit
import json
T_inverted = json.dumps(index, indent=2)
```

- **Result:** { "Atlantis": ["A2", "A7", "F5"], "QuantumMind": ["B2:B4", "F4"], ... }

Effect: Boosts sheet-level compression ratio from $\sim 4.4\times$ to $\sim 14.9\times$ without losing any cell.

3.4 Data-Format-Aware Aggregation

Goal: Further collapse clusters of similarly typed or formatted cells—replacing raw values with their **data type** or **Number Format String (NFS)**.

◆ Implementation Steps

1. Extract NFSs

- Use a library (e.g. **OpenPyXL** or **ClosedXML**) to read the built-in Number Format String for each cell.
- If NFS is missing, **infer** via rule-based matching on the cell's text:

```
python
CopyEdit
RULES = {
    r"^\d{4}-\d{2}-\d{2}$": "Date",
    r"^\d+$": "Integer",
    r"^\d+\.\d+$": "Float",
    r"^\d+%" : "Percentage",
    # ... plus Scientific, Time, Currency, Email, etc.
}
def infer_type(text):
    for pattern, dtype in RULES.items():
        if re.match(pattern, text):
            return dtype
    return "Others"
```

2. Label Each Cell

```
python
CopyEdit
label_map = {}
for addr, cell in Se.items():
    nfs = get_nfs(cell) # library call, or None
    label = nfs or infer_type(cell.value)
    label_map.setdefault(label, []).append(addr)
```

3. Cluster Adjacent Addresses

- As in §3.3, **sort and merge** `label_map[label]` into ranges, but now grouping only those addresses that form **rectangular blocks** of identical label.
- This yields for example:

```
json
```

```
CopyEdit
{
  "IntNum":      ["C2:C4", "C6:C12"],
  "Percentage": ["H5:H6", "H8:H9"],
  ...
}
```

4. Produce Final Aggregated JSON

```
python
CopyEdit
T_aggregated = json.dumps(label_map, indent=2)
```

- **Effect:** Lift compression ratio from $\sim 14.9\times$ up to $\sim 24.8\times$ —now only format/type clusters remain

3.5 Chain of Spreadsheet (CoS)

We wrap our compressor+LLM pipeline into a two-stage procedure:

Stage 1: Table Identification & Boundary Detection

1. Inputs:

- `T_compressed` (the output of §3.4)
- `user_query` (natural-language question)

2. Prompt Template:

```
css
CopyEdit
You are a spreadsheet-understanding assistant.
Given the following compressed sheet representation:
  {T_compressed}
And the query:
  "{user_query}"
Identify the single table that contains the information needed to
answer.
Output exactly:
{
  "table_id": <TABLE_INDEX>,
  "bounds": { "top": r1, "left": c1, "bottom": r2, "right": c2 }
}
```

3. LLM Call & Parsing:

```
python
CopyEdit
detection_resp = llm(prompt_detection)
result = json.loads(detection_resp)
r1, c1, r2, c2 = result["bounds"].values()
```

4. Extract Sub-Table:

Map (r1...r2, c1...c2) back to your compressed or skeleton grid to get `Table_segment`.

Stage 2: Response Generation

1. Prompt Template:

```
css
CopyEdit
You are given a table:
{serialize(Table_segment)}
And the question:
"{user_query}"
Answer concisely, citing cell addresses or formulas where appropriate.
```

2. LLM Call:

```
python
CopyEdit
answer = llm(prompt_answer)
```

3. Output:

Return `answer` to the user.

4 Experiments: Evaluation Setup

4.1 Spreadsheet Table Detection

- **Dataset:**
 - 188 real-world spreadsheets (311 tables) from Dong et al. 2019b, human-validated boundaries.
 - Split by token-count into Small/Medium/Large/Huge.
- **Metric:**
 - **Error-of-Boundary-0 (EoB-0):** exact match of top/left/bottom/right.
 - Report **F1** over all tables.
- **Baselines & Models:**
 - **Baseline:** TableSense-CNN (Dong et al. 2019b).
 - **LLMs:** GPT-4, GPT-3.5; open-source: Llama2, Llama3, Phi3, Mistral-v2.
 - Few-shot prompting or fine-tuning on detection prompts.

4.2 Spreadsheet QA

- **Dataset:**
 - 64 multi-table spreadsheets sampled from the corpus.

- 4–6 handcrafted questions each (search, compare, basic arithmetic), total **307** (Q, A, S) tuples.
- Answers as cell addresses or formulas.
- **Evaluation:**
 - **Exact match** on cell address or formula string.
 - Compare performance when using:
 1. Vanilla encoding + direct QA prompt
 2. Full CoS pipeline

4.2.2 Experiment Setup

1. **Baselines Chosen**
 - **TAPEX** and **Binder** (optimized for single-table QA).
2. **Multi-Table Adaptation**

```
python
CopyEdit
# Step A: Table Detection
table_bounds = detect_table_bounds(Se, question) # from your
fine-tuned detection model

# Step B: Extract & Compress
table_segment = extract_subtable(Se, table_bounds)
compressed = apply_compression(table_segment,
                                use_structural_anchor=True,
                                use_inverted_index=True,
                                use_format_aggregation=True)

# Step C: Format for Baseline
# e.g. for TAPEX: serialize as Markdown table; for Binder: JSON input
schema
baseline_input = format_for_baseline(compressed)
```

3. **Token-Limit Handling**

```
python
CopyEdit
max_tokens = baseline.token_limit() # e.g. 4K tokens

if count_tokens(baseline_input) > max_tokens:
    # Strategy 1: Truncate less-important rows/cols from ends
    baseline_input = truncate_context(baseline_input, max_tokens)
    if count_tokens(baseline_input) > max_tokens:
        # Strategy 2: Split into overlapping chunks
        chunks = split_table(baseline_input, chunk_size=max_tokens,
                              overlap=header_rows)
        # feed each chunk separately and merge answers or pick the
        highest-confidence
```

4. **Answer Evaluation**

- Compare predicted cell address or formula string against ground-truth (Q, A, S) tuples.
 - **Accuracy** = % exact matches.
5. **Model Configuration**
- Use **GPT-4** (instruct-tuned) for your CoS pipeline and for any fine-tuning.
 - Refer to Appendix G for hyperparameters (learning rate, batch size, prompt templates).

4.2.3 Experiment Procedure

1. Detection → QA Loop

```
python
CopyEdit
for (Q, A, S) in spreadsheet_QA_dataset:
    # 1. Table detection (Stage 1 of CoS)
    bounds = detect_table_bounds(S, Q)

    # 2. Extract, compress, and possibly split
    seg = extract_subtable(S, bounds)
    seg_c = apply_compression(seg, modules=[1,2,3])
    seg_chunks = split_if_needed(seg_c, max_tokens)

    # 3. For each chunk: run QA (Stage 2 of CoS)
    answers = [ run_baseline(chunk, Q) for chunk in seg_chunks ]
    pred = merge_or_select(answers)

    # 4. Record correctness
    log_result(Q, A, pred)
```

2. Further Compression & Splitting

- If `modules=[1,2,3]` still exceeds GPT-4’s token limit, deploy the **table-splitting algorithm** (see Appendix M.2):
 - Identify natural “sub-tables” via header detection.
 - Concatenate them into minimal overlapping segments.

3. Logging & Metrics

- Track per-question: detection EoB-0, QA exact-match.
- Aggregate overall and by spreadsheet size (Small/Medium/Large/Huge).

Table 1: Compression Ratios by Module

Modules	Total Tokens	Compression Ratio
None	1,548,577	1.00
1 (Structural Anchor)	350,946	4.41

Modules	Total Tokens	Compression Ratio
2 (Inverted Index)	580,912	2.67
3 (Format Aggregation)	213,890	7.24
1 + 2	103,880	14.91
1 + 3	96,365	16.07
2 + 3	211,445	7.32
1 + 2 + 3	62,469	24.79

Tip: Start with all three modules (1+2+3) for maximal compression; fall back to subsets only if you observe accuracy drops.

3.X Lightweight Heuristics for Structural-Anchor Proposal

Goal

Quickly propose rough table boundaries by purely heuristic rules—so that your LLM can refine them later.

Step 1: Identify High-Discrepancy Rows & Columns

1. **Compute Discrepancy Features** for each row i and column j :
 - **Value Changes:** count how many adjacent cells differ in text or numeric value.
 - **Merged-Cell Presence:** count merged cells in that row/column.
 - **Border Usage:** count cells with non-default borders.
 - **Fill Color:** count cells whose background color \neq default.
 - **Font Emphasis:** count bold/italic cells.
2. **Score & Threshold:**

```
python
CopyEdit
row_score[i] = w1*Δvalue + w2*merges + w3*borders + w4*colors +
w5*fonts
anchors_rows = {i for i in all_rows if row_score[i] ≥ ROW_THRESH}
# similarly for columns → anchors_cols
```

- Choose weights $w_1...w_5$ and `ROW_THRESH` via small grid search on a dev set.

Step 2: Enumerate Candidate Boundaries

Form every rectangle $(r_top, r_bot, c_left, c_right)$ by pairing:

```
python
```

```
CopyEdit
candidates = []
for r1 in anchors_rows:
    for r2 in anchors_rows:
        if r2 <= r1: continue
        for c1 in anchors_cols:
            for c2 in anchors_cols:
                if c2 <= c1: continue
                candidates.append((r1, r2, c1, c2))
```

Step 3: Filter Unreasonable Candidates

For each $(r1, r2, c1, c2)$:

1. **Size Check:** require $(r2-r1+1) \geq \text{MIN_ROWS}$ and $(c2-c1+1) \geq \text{MIN_COLS}$.
2. **Density Check:**

```
python
CopyEdit
area = (r2-r1+1)*(c2-c1+1)
non_empty = count_nonempty_cells(r1, r2, c1, c2)
sparsity = non_empty / area
```

– Reject if $\text{sparsity} < \text{SPARSITY_MIN}$ or $> \text{SPARSITY_MAX}$.

3. **Header-Row/Col Evidence:**
 - Check top row $r1$ has high text-ratio ($\geq \text{HEADER_TEXT_RATIO}$).
 - Check left col $c1$ has $\geq \text{HEADER_LABEL_RATIO}$ of non-numeric values. – Reject if neither edge looks header-like.
-

Step 4: Resolve Overlapping Candidates

1. **Detect Overlaps:**

Two boundaries overlap if their intersection area is non-empty.
 2. **Pairwise Pruning:**

For each overlapping pair A and B:

 - If $|A.r1 - B.r1| \leq \text{DELTA_R}$ (close top edges), compare:
 - **Header Score** = $\text{text_ratio}(\text{top row}) + \text{format_ratio}(\text{year/date patterns})$.
 - **Keep** the one with higher header score; **discard** the other.
 3. **Repeat** until no overlaps remain.
-

Step 5: Expand to Structural Anchors

From the final set of boundaries, collect:

```
python
CopyEdit
raw_anchor_rows = {r1, r2 for each boundary}
raw_anchor_cols = {c1, c2 for each boundary}
```

Then **include neighbors within k units**:

```
python
CopyEdit
k = 4 # preserves >97% true boundary lines
anchors_rows = U_{p ∈ raw_anchor_rows} {i | |i-p| ≤ k}
anchors_cols = U_{q ∈ raw_anchor_cols} {j | |j-q| ≤ k}
```

These expanded sets become your structural anchors $A = \{rp, cq\}$ for §3.2.

Implementation Tips

- **Parameter Tuning:**
 - Calibrate `ROW_THRESH`, `MIN_ROWS`, `SPARSITY_MIN`, etc. on a small validation set.
- **Efficiency:**
 - Precompute per-row/col scores once.
 - Discard obviously tiny candidates early to avoid combinatorial explosion.
- **Robustness:**
 - If boundaries are sparse, falling back to a higher k ensures you don't lose headers/notes.

F Spreadsheet Table Detection Test Dataset Partition

1. Encode Each Sheet in Markdown-Like Style (per §3.1)

- **Text Lines:**
For every non-empty cell (i, j) produce

```
less
CopyEdit
|{col_letter(j)}{i},{cell.value}|
```

- **Format Lines:**
In parallel, for each cell output

```
less
CopyEdit
|{col_letter(j)}{i},{comma-joined list of format attributes}|
```

- **Concatenate:**
Join all text lines into one big string, and likewise for format lines (or interleave them, depending on your tokenizer setup).
-

2. Count Tokens

- **Choose Tokenizer:**
Use the same tokenizer your detection model (e.g. GPT-4) expects—for instance, OpenAI's `tiktoken`.
- **Compute:**

```
python
CopyEdit
import tiktoken
enc = tiktoken.get_encoding("cl100k_base")
tokens = enc.encode(full_markdown_string)
token_count = len(tokens)
```

3. Assign Size Buckets

Bucket	Token Range
Small	$< 4\,000$
Medium	$4\,000 \leq \text{tokens} < 8\,000$
Large	$8\,000 \leq \text{tokens} < 32\,000$
Huge	$\text{tokens} \geq 32\,000$

```
python
CopyEdit
if token_count < 4_000:
    bucket = "Small"
elif token_count < 8_000:
    bucket = "Medium"
elif token_count < 32_000:
    bucket = "Large"
else:
    bucket = "Huge"
```

- **Store** this `bucket` label alongside each sheet in your test metadata.
-

4. Example Encoding Snippet

Text Input

```
less
```

```
CopyEdit
|B2,Table 4: Diesel-driven passenger cars, 2015|
|C2,| |D2,| |E2,| |F2,| |G2,| |H2,| | |
|B3,| |C3,| |D3,| |E3,| |F3,| |G3,| |H3,|
|B4,| |C4,| |D4,| |E4,| |F4,| |G4,| |H4,|
|B5,| |C5,Diesel engine|D5,| |E5,| |F5,Share of all passenger cars
(%)|G5,| |H5,|
...
```

Format Input

```
mathematica
CopyEdit
|B2,Font Bold|C2,|D2,|E2,|F2,|G2,|H2,|
|B3,|C3,|D3,|E3,|F3,|G3,|H3,|
|B4,Bottom Border|C4,Bottom Border|D4,Bottom Border|E4,Bottom
Border|F4,Bottom Border|G4,Bottom Border|H4,Bottom Border|
|B5,Top Border,Right Border,Fill Color,Font Bold|
C5,Top Border,Bottom Border,Left Border,Fill Color,Font Bold|
D5,Top Border,Bottom Border,Fill Color,Font Bold|
E5,Top Border,Bottom Border,Right Border,Fill Color,Font Bold|
F5,Top Border,Bottom Border,Left Border,Fill Color,Font Bold|
G5,Top Border,Bottom Border,Fill Color,Font Bold|
H5,Top Border,Bottom Border,Fill Color,Font Bold|
```

1. Example Spreadsheet QA Data Item

Each QA example is a tuple (question, ground_truth, prompt):

```
python
CopyEdit
qa_item = {
    "question": "What were the highest temperatures in Washington DC in
1998?",
    "ground_truth": ["X23", "X24"],
    "prompt": build_prompt(
        instruction=YOUR_INSTRUCTION_STRING,
        sheet_encoding=T_aggregated # the fully compressed JSON/Markdown
    )
}
```

- **question:** the user’s natural-language query.
- **ground_truth:** list of cell addresses or formula(s) expected.
- **prompt:** your “Instruction + Encoded Spreadsheet” string fed into the LLM.

2. Cost Calculation

Track token usage and API cost like this:

```
python
CopyEdit
# Constants (per-1K-token prices)
PRICE_GPT35 = 0.0005    # USD per 1K tokens
PRICE_GPT4   = 0.03

# Measured average compressed tokens per sheet
avg_tokens = 62_000 / 198 # ~313 tokens/sheet

# Cost per task
cost_gpt35 = avg_tokens * PRICE_GPT35 / 1000
cost_gpt4   = avg_tokens * PRICE_GPT4   / 1000

print(f"GPT-3.5 turbo cost: ${cost_gpt35:.6f}") # ~$0.000157
print(f"GPT-4 cost: ${cost_gpt4:.6f}") # ~$0.00939
```

- **Baseline (vanilla)** uses ~1,548,000 tokens → ~\$0.00391 (GPT-3.5) / \$0.235 (GPT-4).
 - **Compressed** uses ~62,000 tokens → saves ~96% in cost.
-

3. Qualitative Comparison with TableSense-CNN

Example from **Figure 10**:

```
python
CopyEdit
baseline_regions = [("A1", "G44"), ("K5", "M14"), ("K16", "M38"),
("Q20", "W29")]
sllm_regions     = [("A1", "G44"), ("K5", "R14"), # extended K5→R14
("K16", "M38"), ("Q20", "W29")]

# Compute what SLLM added:
added = set(sllm_regions) - set(baseline_regions)
# added == {"K5", "R14"}
```

- **Insight:** SLLM correctly includes column **R5:R14**, which TableSense-CNN missed because it's semantically linked (percentages of earlier columns) but spatially distant.

L.1 Vanilla Prompt Template for Table Detection

```
text
CopyEdit
INSTRUCTION:
You are given a serialized spreadsheet as a single line of cell tuples in
row-major order.
Each tuple is formatted as `<Address>,<Value>`, and tuples are separated by
`|`.
Empty cells appear as `<Address>`,``.
Identify every table in this sheet (header + data rows only; exclude titles
or comments).
Return a JSON array of range strings, e.g.:
```



```
["A2:D5", "K5:M14"]
```

Do ****not**** output any other text or explanation.

INPUT:

[Encoded Spreadsheet]

L.2 Compressed-Input Prompt Template for Table Detection

text

CopyEdit

INSTRUCTION:

You are given a compressed spreadsheet encoding as a JSON object mapping each cell text or format label to one or more addresses/ranges, e.g.:

```
{
  "Year": ["A1"],
  "IntNum": ["B2:B10"],
  "#,##0": ["C2:C10"],
  ...
}
```

Detect all tables (header + data) in the sheet.

Return a JSON array of range strings, e.g.:

```
["A1:F9", "K5:R14"]
```

Do ****not**** include titles or comments, and do ****not**** add any other text.

INPUT:

[Compressed JSON]

L.3 Prompt Templates for Spreadsheet QA (CoS)

Stage 1 – Table Identification

text

CopyEdit

INSTRUCTION:

You are given:

- A compressed spreadsheet (JSON as above).
- A question about this sheet.

Determine which single table contains the answer.

Return exactly one JSON object:

```
{"range": "A1:F9"}
```

Do ****not**** add any other text or explanation.

INPUT:

Spreadsheet: [Compressed JSON]
Question: "<Your Question>"

Stage 2 – Answer Generation

```
text
CopyEdit
INSTRUCTION:
You are given:
  • A single table (Markdown-style or JSON).
  • A question whose answer lies within this table.

Locate the answer and return exactly one JSON object:

{"answer": "<CellAddress or Formula>"}

Examples: `{"answer": "B3"}`, `{"answer": "SUM(A2:A10)"}\`.
Do not add any other text or explanation.

INPUT:
Table: [Table Representation]
Question: "<Your Question>"
```

Tips for Use in Your Guide:

- Embed these templates verbatim in your code or README.
- Fill in [Encoded Spreadsheet], [Compressed JSON], and <Your Question> programmatically.
- Always parse the LLM's output as strict JSON to avoid downstream errors.

M.1 Identical Cell Aggregation (Algorithm 1)

Input:

- $nfs[m][n]$ — 2D matrix of Number Format Strings or inferred type labels for each cell.

Output:

- $areas$ — list of tuples $((r1, c1), (r2, c2), val_type)$ denoting each rectangular block of identical type.

```
pseudo
CopyEdit
1. Let  $m, n$  = dimensions of  $nfs$ 
2.  $visited \leftarrow$  2D array  $[m][n]$ , all False
3.  $areas \leftarrow []$ 
4.  $FormatDict \leftarrow$  mapping from NFS string  $\rightarrow$  canonical  $val\_type$ 

5. define  $dfs(r, c, val\_type)$ :
6.   if  $visited[r][c]$  OR  $val\_type \neq FormatDict[nfs[r][c]]$ :
```

```
7.     return bounds (r,c,r,c)  # single-cell block
8.     visited[r][c] ← True
9.     bounds ← (r, c, r, c)
10.    for each neighbor (tr, tc) in {up, down, left, right}:
11.        if not visited[tr][tc] AND val_type == FormatDict[nfs[tr][tc]]:
12.            sub_bounds = dfs(tr, tc, val_type)
13.            expand bounds to include sub_bounds
14.    return bounds

15. for r in 0 ... m-1:
16.     for c in 0 ... n-1:
17.         if not visited[r][c]:
18.             val_type = FormatDict[nfs[r][c]]
19.             (r1,c1,r2,c2) = dfs(r, c, val_type)
20.             areas.append(((r1,c1), (r2,c2), val_type))

21. return areas
```

M.2 Table-Split QA Algorithm (Algorithm 2)

Input:

- question — string
- region — 2D table (list of rows), already compressed or extracted

Output:

- answers — list of answer strings (cell address or formula)

pseudo

CopyEdit

```
1. headers ← []
2. answers ← []

3. if calculateTokens(region) ≤ 4096:
4.     return [ answer_question(question, region) ]

5. else:
6.     headers = predict_header(region)
7.     body = region[ len(headers) : ]  # rows below header

8.     for i in 0 ... len(body)-1:
9.         # build a small sub-table: header + 3 rows at a time
10.        sub_table = headers + body[i : i+3]
11.        answer = answer_question(question, sub_table)
12.        answers.append(answer)

13. return answers
```

- **calculateTokens**: counts tokens in your chosen encoding.
- **predict_header**: identifies header rows (e.g. via heuristic or model).
- **answer_question**: runs your Stage 2 QA LLM prompt on the given sub-table.

